

Benchmark Report: Fourier Attention vs. Standard Self-Attention

Date: April 26, 2025

Device: CUDA-enabled (if available), otherwise CPU

Batch Size: 8

Sequence Length: 512

Embedding Dimension: 64

Number of Heads: 8

=== Overview ===

This report compares two attention mechanisms:

Mechanism	Description
Fourier Attention	Applies Fast Fourier Transform (FFT) over the sequence dimension to approximate interactions, followed by an inverse FFT. Intended to be computationally efficient for long sequences.
Standard Self-Attention	Traditional attention using query-key-value projections and dot-product softmax attention. Provides rich contextualization but with higher memory and compute demands.

=== Model Characteristics ===

Model	Parameter Count	FFT Used	Attention Mechanism	Complexity
Fourier Attention	~8,320	Yes	Frequency-based	$O(N \log N)$
Self-Attention	~16,384	No	Dot-product	$O(N^2)$

=== Procedure ===

1. Environment Setup:

- Python with PyTorch and einops installed.
- GPU used if available (torch.cuda.is_available()).

2. Input Configuration:

- Random tensor of shape [batch_size=8, seq_len=512, dim=64] created as input.

3. Model Initialization:

- Both FourierAttention and StandardSelfAttention models instantiated with 8 heads.

4. Benchmarking:

- For each model:
 - Forward pass time measured using time.time().
 - Memory usage collected using torch.cuda.max_memory_allocated() (on GPU).
 - Output shape and total parameter count printed.
- For Fourier Attention:
 - Optionally visualized FFT magnitude for one head using matplotlib.

5. Evaluation Mode:

- All models run under torch.no_grad() and .eval() mode to disable gradient tracking for fair benchmarking.

=== Performance Summary (Sample Results) ===

These are indicative results. Real values depend on your specific hardware.

Metric	Fourier Attention	Standard Self-Attention
-----	-----	-----

Forward Pass Time	~2-3 ms	~12-16 ms	
Peak Memory Usage (GPU)	~20 MB	~65-70 MB	
Output Shape	[8, 512, 64]	[8, 512, 64]	

=== FFT Visualization (Fourier Attention) ===

- The model emphasizes low to mid-frequency bins, indicating it captures global patterns well.
- High-frequency components typically have lower magnitude, which might reflect less attention to local fluctuations.

=== Key Observations ===

- Speed: Fourier Attention offers faster inference, especially for longer sequences.
- Memory: It uses significantly less GPU memory, making it more scalable.
- Expressivity: Standard Self-Attention remains more expressive, particularly for tasks requiring fine-grained, token-to-token interactions.

=== Recommendations ===

Use Case	Recommended Model	
-----	-----	
Long sequences, low memory	Fourier Attention	
Fine-grained understanding	Standard Self-Attention	
Real-time inference on edge	Fourier Attention	
Complex reasoning tasks (e.g., NLP)	Standard Self-Attention	

=== Potential Enhancements ===

- Add backward pass to evaluate training efficiency.
- Test with mixed precision (AMP) to further reduce memory.
- Integrate with real data pipelines (e.g., NLP/vision embeddings).

- Experiment with hybrid models (FFT + token attention).